

LEARNING AS OPTIMIZATION: MOTIVATION

Learning as optimization: general procedure

- Goal: Learn parameter θ (or weight vector \mathbf{w})
- Dataset: $D = \{(x_1, y_1), \dots, (x_n, y_n)\}$
- Write down loss function: how well \mathbf{w} fits the data D as a function of \mathbf{w}
 - Common choice: $\log \Pr(D|\mathbf{w})$
- Maximize by differentiating
 - Then **gradient descent**: repeatedly take a small step in the direction of the gradient

Learning as optimization: general procedure for SGD (stochastic gradient descent)

- **Big-data** problem: we *don't* want to load all the data D into memory, and the gradient depends on all the data
- **Solution:**
 - pick a small subset of examples $B \ll D$
 - **approximate** the gradient using them
 - “on average” this is the right direction
 - take a step in that direction
 - repeat....
- Math: find gradient of w for a *single* example, not a dataset

B = one
example is
a very
popular
choice

SGD vs streaming

- Streaming:
 - pass through the data *once*
 - hold model + one example in memory
 - update model for each example
- Stochastic gradient:
 - pass through the data *multiple times*
 - stream through a disk file repeatedly
 - hold model + B examples in memory
 - update model *via gradient step*

B = **one** example is a very popular choice

its simple 😊

sometimes its cheaper to evaluate 100 examples at once than one example 100 times 😞



Efficient Logistic Regression with Stochastic Gradient Descent

William Cohen

Learning as optimization for logistic regression

- Goal: Learn the parameter \mathbf{w} of the classifier

$$P(Y = y|X = \mathbf{x}, \mathbf{w}) = \frac{1}{1 + e^{-\mathbf{x} \cdot \mathbf{w}}}$$

- Probability of a single example $P(y|\mathbf{x}, \mathbf{w})$ would be

$$P(Y = y|X = \mathbf{x}, \mathbf{w}) = \begin{cases} \frac{1}{1+e^{-\mathbf{x} \cdot \mathbf{w}}} & \text{if } y = 1 \\ 1 - \frac{1}{1+e^{-\mathbf{x} \cdot \mathbf{w}}} & \text{if } y = 0 \end{cases}$$

- Or with logs:

$$\log P(Y = y|X = \mathbf{x}, \mathbf{w}) = \begin{cases} \log p & \text{if } y = 1 \\ \log(1 - p) & \text{if } y = 0 \end{cases}$$

$$p \equiv \frac{1}{1 + e^{-\mathbf{x} \cdot \mathbf{w}}} = \frac{1}{1 + \exp(-\sum_j x^j w^j)}$$

$$\log P(Y = y|X = \mathbf{x}, \mathbf{w}) = \begin{cases} \log p & \text{if } y = 1 \\ \log(1 - p) & \text{if } y = 0 \end{cases}$$

$$\frac{\partial}{\partial w^j} \log P(Y = y|X = \mathbf{x}, \mathbf{w}) = \begin{cases} \frac{1}{p} \frac{\partial}{\partial w^j} p & \text{if } y = 1 \\ \frac{1}{1-p} \left(-\frac{\partial}{\partial w^j} p \right) & \text{if } y = 0 \end{cases}$$

$$(\log f)' = \frac{1}{f} f'$$

$$p \equiv \frac{1}{1 + e^{-\mathbf{x} \cdot \mathbf{w}}} = \frac{1}{1 + \exp(-\sum_j x^j w^j)}$$

$$p \equiv \frac{1}{1 + e^{-\mathbf{x} \cdot \mathbf{w}}} = \frac{1}{1 + \exp(-\sum_j x^j w^j)}$$

$$1 - p = \frac{1 + \exp(-\sum_j x^j w^j)}{1 + \exp(-\sum_j x^j w^j)} - \frac{1}{1 + \exp(-\sum_j x^j w^j)} = \boxed{\frac{\exp(-\sum_j x^j w^j)}{1 + \exp(-\sum_j x^j w^j)}}$$

$$\boxed{\frac{\partial}{\partial w^j} p} = \frac{\partial}{\partial w^j} (1 + \exp(-\sum_j x^j w^j))^{-1} \quad (e^f)' = e^f f'$$

$$= (-1)(1 + \exp(-\sum_j x^j w^j))^{-2} \frac{\partial}{\partial w^j} \exp(-\sum_j x^j w^j)$$

$$= (-1)(1 + \exp(-\sum_j x^j w^j))^{-2} \exp(-\sum_j x^j w^j) (-x^j)$$

$$\stackrel{\text{P}}{=} \boxed{\frac{1}{1 + \exp(-\sum_j x^j w^j)}} \boxed{\frac{\exp(-\sum_j x^j w^j)}{1 + \exp(-\sum_j x^j w^j)}} x^j$$

$$\frac{\partial}{\partial w^j} p = p(1 - p)x^j$$

$$\log P(Y = y|X = \mathbf{x}, \mathbf{w}) = \begin{cases} \log p & \text{if } y = 1 \\ \log(1 - p) & \text{if } y = 0 \end{cases}$$

$$\frac{\partial}{\partial w^j} \log P(Y = y|X = \mathbf{x}, \mathbf{w}) = \begin{cases} \frac{1}{p} \frac{\partial}{\partial w^j} p & \text{if } y = 1 \\ \frac{1}{1-p} \left(-\frac{\partial}{\partial w^j} p \right) & \text{if } y = 0 \end{cases}$$

$$\frac{\partial}{\partial w^j} p = p(1 - p)x^j$$

$$\frac{\partial}{\partial w^j} \log P(Y = y|X = \mathbf{x}, \mathbf{w}) = \begin{cases} \frac{1}{p} p(1 - p)x^j = (1 - p)x^j & \text{if } y = 1 \\ \frac{1}{1-p} (-1)p(1 - p)x^j = -px^j & \text{if } y = 0 \end{cases}$$

$$\frac{\partial}{\partial w^j} \log P(Y = y|X = \mathbf{x}, \mathbf{w}) = (y - p)x^j$$

$$\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} + \lambda(y - p)\mathbf{x}$$

$$\log P(Y = y|X = \mathbf{x}, \mathbf{w}) = \begin{cases} \log p & \text{if } y = 1 \\ \log(1 - p) & \text{if } y = 0 \end{cases}$$

$$p = \sigma(\mathbf{x} \cdot \mathbf{w})$$

Magically, when we differentiate, we end up with something very simple and elegant.....

$$\frac{\partial}{\partial \mathbf{w}} L(\mathbf{w} | y, \mathbf{x}) = (y - p)\mathbf{x}$$

$$\frac{\partial}{\partial w^j} L(\mathbf{w} | y, \mathbf{x}) = (y - p)x^j$$

The update for gradient descent is just:

$$\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} + \lambda(y - p)\mathbf{x}$$

**Logistic regression has a
sparse update**

An observation: sparsity!

$$\frac{\partial}{\partial w^j} \log P(Y = y | X = \mathbf{x}, \mathbf{w}) = (y - p)x^j$$

Key computational point:

- if $x^j = 0$ then the gradient of w^j is zero
- so when processing an example you only need to update weights for the **non-zero** features of an example.

Learning as optimization for logistic regression

- The algorithm: $\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} + \lambda(y - p)\mathbf{x}$

1. Initialize a hashtable W

2. For $t = 1, \dots, T$

- For each example \mathbf{x}_i, y_i :
 - *do this in random order*
 - Compute the prediction for \mathbf{x}_i :

$$p_i = \frac{1}{1 + \exp(-\sum_{j: x_i^j > 0} x_i^j w^j)}$$

- For each non-zero feature of \mathbf{x}_i with index j and value x^j :
 - * If j is not in W , set $W[j] = 0$.
 - * Set $W[j] = W[j] + \lambda(y_i - p_i)x^j$

3. Output the hash table W .

REGULARIZED LOGISTIC REGRESSION

Regularized logistic regression

- Replace LCL

$$\log P(Y = y|X = \mathbf{x}, \mathbf{w}) = \begin{cases} \log p & \text{if } y = 1 \\ \log(1 - p) & \text{if } y = 0 \end{cases}$$

- with LCL + penalty for large weights, eg

$$LCL - \mu \sum_{j=1}^d (w^j)^2$$

- So:

$$\frac{\partial}{\partial w^j} \log P(Y = y|X = \mathbf{x}, \mathbf{w}) = (y - p)x^j$$

- becomes:

$$\frac{\partial}{\partial w^j} \log P(Y = y|X = \mathbf{x}, \mathbf{w}) - \mu \sum_{j=1}^d (w^j)^2 = (y - p)x^j - 2\mu w^j$$

Regularized logistic regression

- Replace LCL

$$\log P(Y = y|X = \mathbf{x}, \mathbf{w}) = \begin{cases} \log p & \text{if } y = 1 \\ \log(1 - p) & \text{if } y = 0 \end{cases}$$

- with LCL + penalty for large weights, eg

$$LCL - \mu \sum_{j=1}^d (w^j)^2$$

- So the update for w_j becomes:

$$w^j = w^j + \lambda((y - p)x^j - 2\mu w^j)$$

- Or

$$w^j = w^j + \lambda(y - p)x^j - \lambda 2\mu w^j$$

Learning as optimization for logistic regression

• Algorithm: $\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} + \lambda(y - p)\mathbf{x}$

1. Initialize a hashtable W

2. For $t = 1, \dots, T$

- For each example \mathbf{x}_i, y_i :
 - *do this in random order*
 - Compute the prediction for \mathbf{x}_i :

$$p_i = \frac{1}{1 + \exp(-\sum_{j: x_i^j > 0} x_i^j w^j)}$$

- For each non-zero feature of \mathbf{x}_i with index j and value x^j :
 - * If j is not in W , set $W[j] = 0$.
 - * Set $W[j] = W[j] + \lambda(y_i - p_i)x^j$

3. Output the hash table W .

Learning as optimization for regularized logistic regression

- Algorithm: $w^j = w^j + \lambda(y - p)x^j - \lambda 2\mu w^j$

1. Initialize a hashtable W

2. For $t = 1, \dots, T$

- For each example \mathbf{x}_i, y_i :
 - Compute the prediction for \mathbf{x}_i :

$$p_i = \frac{1}{1 + \exp(-\sum_{j: x_i^j > 0} x_i^j w^j)}$$

- For each ~~non-zero~~ feature of \mathbf{x}_i with index j and value x^j :
 - If j is not in W , set $W[j] = 0$.
 - Set $W[j] = W[j] + \lambda(y - p)x^j - \lambda 2\mu w^j$

3. Output the hash table W .

Time goes from $O(nT)$ to $O(mVT)$ where

- n = number of non-zero entries,
- m = number of examples
- V = number of features
- T = number of passes over data

This change is very important for large datasets

- We've lost the ability to do *sparse* updates
- This makes learning *much much* more expensive
 - 2×10^6 examples
 - 2×10^8 non-zero entries
 - 2×10^6 + features
 - 10,000x slower (!)

Time goes from $O(nT)$ to $O(mVT)$
where

- n = number of non-zero entries,
- m = number of examples
- V = number of features
- T = number of passes over data

SPARSE UPDATES FOR REGULARIZED LOGISTIC REGRESSION

Learning as optimization for regularized logistic regression

- Final algorithm: $w^j = w^j + \lambda(y - p)x^j - \lambda 2\mu w^j$
- Initialize hashtable W
- For each iteration $t=1, \dots, T$
 - For each example (\mathbf{x}_i, y_i)
 - $p_i = \dots$
 - For each feature $W[j]$
 - $W[j] = W[j] - \lambda 2\mu W[j]$
 - If $x_i^j > 0$ then
 - » $W[j] = W[j] + \lambda(y_i - p^i)x_j$

Learning as optimization for regularized logistic regression

- Final algorithm: $w^j = w^j + \lambda(y - p)x^j - \lambda 2\mu w^j$
- Initialize hashtable W
- For each iteration $t=1, \dots, T$
 - For each example (\mathbf{x}_i, y_i)
 - $p_i = \dots$
 - For each feature $W[j]$
 - $W[j] *= (1 - \lambda 2\mu)$
 - If $x_i^j > 0$ then
 - » $W[j] = W[j] + \lambda(y_i - p^i)x_j$

Learning as optimization for regularized logistic regression

- Final algorithm: $w^j = w^j + \lambda(y - p)x^j - \lambda 2\mu w^j$
- Initialize hashtable W
- For each iteration $t=1,\dots,T$
 - For each example (\mathbf{x}_i, y_i)
 - $p_i = \dots$
 - For each feature $W[j]$
 - If $x_i^j > 0$ then
 - » $W[j] *= (1 - \lambda 2\mu)^A$
 - » $W[j] = W[j] + \lambda(y_i - p^i)x_j$

A is number of examples seen since the last time we did an **x>0** update on $W[j]$

Learning as optimization for regularized logistic regression

- Final algorithm: $w^j = w^j + \lambda(y - p)x^j - \lambda 2\mu w^j$
- Initialize hashtables W, A and set $k=0$
- For each iteration $t=1, \dots, T$
 - For each example (\mathbf{x}_i, y_i)
 - $p_i = \dots$; $k++$
 - For each feature $W[j]$
 - If $x_i^j > 0$ then
 - » $W[j] *= (1 - \lambda 2\mu)^{k-A[j]}$
 - » $W[j] = W[j] + \lambda(y_i - p^i)x_j$
 - » $A[j] = k$

$k-A[j]$ is number of examples seen since the last time we did an $x > 0$ update on $W[j]$

Learning as optimization for regularized logistic regression

- Final algorithm: $w^j = w^j + \lambda(y - p)x^j - \lambda 2\mu w^j$
- Initialize hashtables W, A and set $k=0$
- For each iteration $t=1, \dots, T$
 - For each example (\mathbf{x}_i, y_i)
 - $p_i = \dots; k++$
 - For each feature $W[j]$
 - If $x_i^j > 0$ then
 - » $W[j] *= (1 - \lambda 2\mu)^{k-A[j]}$
 - » $W[j] = W[j] + \lambda(y_i - p^i)x_j$
 - » $A[j] = k$

- k = “clock” reading
- $A[j]$ = clock reading last time feature j was “active”
- we implement the “weight decay” update using a “lazy” strategy: weights are decayed in one shot when a feature is “active”

Learning as optimization for regularized logistic regression

- Final algorithm: $w^j = w^j + \lambda(y - p)x^j - \lambda 2\mu w^j$
- Initialize hashtables W, A and set $k=0$
- For each iteration $t=1, \dots, T$
 - For each example (x_i, y_i)
 - $p_i = \dots; k++$
 - For each feature $W[j]$
 - If $x_i^j > 0$ then
 - » $W[j] *= (1 - \lambda 2\mu)^{k-A[j]}$
 - » $W[j] = W[j] + \lambda(y_i - p^i)x_j$
 - » $A[j] = k$

Time goes from $O(nT)$ to $O(mVT)$ where

- n = number of non-zero entries,
- m = number of examples
- V = number of features
- T = number of passes over data

Memory use doubles.

Comments

- What's happened here:
 - Our update involves a *sparse part* and a *dense part*
 - Sparse: empirical loss on this example
 - Dense: regularization loss – not affected by the example
 - We remove the *dense part* of the update
 - Old example update:
 - for each feature { do something example-independent}
 - For each active feature { do something example-dependent}
 - New example update:
 - For each active feature :
 - » {simulate the prior example-independent updates}
 - » {do something example-dependent}